

hw3

Tinglei Wu

2/14/2022

```
library('splines')      ## for 'bs'
library('dplyr')        ## for 'select', 'filter', and others
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library('magrittr')     ## for '%<>%' operator
library('glmnet')       ## for 'glmnet'
```


```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-3
```

```
### Linear regression examples ###
```

```
## load prostate data
prostate <-
  read.table(url(
    'https://web.stanford.edu/~hastie/ElemStatLearn/datasets/prostate.data'))

pairs(prostate)
```



```
## split prostate into testing and training subsets
```

```
prostate_train <- prostate %>%
  filter(train == TRUE) %>%
  select(-train)
```

```
summary(prostate_train)
```

```
##      lcavol      lweight      age      lbph
## Min.   :-1.3471  Min.   :2.375  Min.   :41.00  Min.   :-1.38629
## 1st Qu.: 0.4883  1st Qu.:3.330  1st Qu.:61.00  1st Qu.: -1.38629
## Median : 1.4679  Median :3.599  Median :65.00  Median :-0.05129
## Mean   : 1.3135  Mean   :3.626  Mean   :64.75  Mean   : 0.07144
## 3rd Qu.: 2.3491  3rd Qu.:3.884  3rd Qu.:69.00  3rd Qu.: 1.54751
## Max.    : 3.8210  Max.    :4.780  Max.    :79.00  Max.    : 2.32630
##      svi      lcp      gleason      pgg45
## Min.   :0.0000  Min.   :-1.3863  Min.   :6.000  Min.   : 0.00
## 1st Qu.:0.0000  1st Qu.: -1.3863  1st Qu.:6.000  1st Qu.: 0.00
## Median :0.0000  Median :-0.7985  Median :7.000  Median : 15.00
## Mean   :0.2239  Mean   :-0.2142  Mean   :6.731  Mean   : 26.27
## 3rd Qu.:0.0000  3rd Qu.: 0.9948  3rd Qu.:7.000  3rd Qu.: 50.00
## Max.    :1.0000  Max.    : 2.6568  Max.    :9.000  Max.    :100.00
##      lpsa
## Min.   :-0.4308
## 1st Qu.: 1.6673
## Median : 2.5688
## Mean   : 2.4523
## 3rd Qu.: 3.3652
## Max.    : 5.4775
```

```
prostate_test <- prostate %>%
  filter(train == FALSE) %>%
  select(-train)
```

##Question 2:

```
cor(prostate_train)
```

```
##          lcavol    lweight      age      lbph      svi      lcp
## lcavol  1.00000000 0.30023199 0.2863243 0.06316772 0.5929491 0.69204308
## lweight 0.30023199 1.00000000 0.3167235 0.43704154 0.1810545 0.15682859
## age     0.28632427 0.31672347 1.0000000 0.28734645 0.1289023 0.17295140
## lbph    0.06316772 0.43704154 0.2873464 1.00000000 -0.1391468 -0.08853456
## svi     0.59294913 0.18105448 0.1289023 -0.13914680 1.0000000 0.67124021
## lcp     0.69204308 0.15682859 0.1729514 -0.08853456 0.6712402 1.00000000
## gleason 0.42641407 0.02355821 0.3659151 0.03299215 0.3068754 0.47643684
## pgg45   0.48316136 0.07416632 0.2758057 -0.03040382 0.4813577 0.66253335
## lpsa    0.73315515 0.48521519 0.2276424 0.26293763 0.5568864 0.48920320
##          gleason    pgg45    lpsa
## lcavol  0.42641407 0.48316136 0.7331551
## lweight 0.02355821 0.07416632 0.4852152
## age     0.36591512 0.27580573 0.2276424
## lbph    0.03299215 -0.03040382 0.2629376
## svi     0.30687537 0.48135774 0.5568864
## lcp     0.47643684 0.66253335 0.4892032
## gleason 1.00000000 0.75705650 0.3424278
## pgg45   0.75705650 1.00000000 0.4480480
## lpsa    0.34242781 0.44804795 1.0000000
```

##Question 3,4:

```
lcavol_out<- lm(lcavol ~ ., data=prostate_train)
lcavol_out
```

```
##
## Call:
## lm(formula = lcavol ~ ., data = prostate_train)
##
## Coefficients:
## (Intercept)      lweight      age      lbph      svi      lcp
##   -2.173357   -0.113370   0.020102  -0.056981   0.035116   0.418455
##      gleason      pgg45      lpsa
##    0.224387   -0.009113   0.575455
```

##Question 5:

```
L2_loss <- function(y, yhat)
  (y-yhat)^2
err <- function(dat, fit, loss=L2_loss)
  mean(loss(dat$lcavol, predict(fit, newdata=dat)))
## training error
err(prostate_train, lcavol_out)
```

```
## [1] 0.4383709
```

```
## testing error
err(prostate_test, lcavol_out)
```

```
## [1] 0.5084068
```

##Question 6:

```
form <- lcavol ~ lweight + age + lbph + lcp + pgg45 + lpsa + svi + gleason
x_inp <- model.matrix(form, data=prostate_train)
y_out <- prostate_train$lcavol
glmnet <- glmnet(x=x_inp, y=y_out, lambda=seq(0.8, 0, -0.05), alpha = 0)
print(glmnet$beta)
```

```
## 9 x 17 sparse Matrix of class "dgCMatrix"
```

```
##      [[ suppressing 17 column names 's0', 's1', 's2' ... ]]
```

```
##
## (Intercept)      .              .              .              .              .
## lweight      0.126233723  0.123269298  0.119863278  0.1158716819  0.1111931576
## age          0.011482623  0.011670731  0.011877123  0.0121010799  0.0123455384
## lbph        -0.007449919 -0.008367575 -0.009363254 -0.0104653587 -0.0116905450
## lcp          0.207036526  0.211463094  0.216338206  0.2216006088  0.2273130852
## pgg45        0.001473672  0.001329880  0.001167084  0.0009827841  0.0007732439
## lpsa         0.290917892  0.298034308  0.305629735  0.3138070249  0.3226494596
## svi         0.394617510  0.393975652  0.392543932  0.3903246504  0.3871479640
## gleason      0.132832538  0.133493765  0.134175184  0.1349324324  0.1357928799
##
## (Intercept)      .              .              .              .              .
## lweight      0.1057367423  0.0992845548  9.167569e-02  0.0826683974
## age          0.0126162948  0.0129137196  1.324527e-02  0.0136179442
## lbph        -0.0130587545 -0.0145982826 -1.633988e-02 -0.0183250234
## lcp          0.2336122365  0.2404842462  2.481056e-01  0.2566365841
## pgg45        0.0005320808  0.0002563424 -6.281817e-05 -0.0004352185
## lpsa         0.3322362737  0.3427336970  3.542772e-01  0.3670625821
## svi         0.3827345083  0.3769402232  3.693573e-01  0.3595330231
## gleason      0.1368039677  0.1380061651  1.394762e-01  0.1413173664
##
## (Intercept)      .              .              .              .              .
## lweight      0.0719539790  0.059127616  0.043652593  0.024777969  0.001504802
## age          0.0140406458  0.014526957  0.015088490  0.015748487  0.016532948
## lbph        -0.0206075810 -0.023258103 -0.026377963 -0.030098852 -0.034621150
## lcp          0.2662889676  0.277447149  0.290342311  0.305728439  0.324372008
## pgg45       -0.0008738898 -0.001398912 -0.002031353 -0.002810371 -0.003788173
## lpsa         0.3813402190  0.397429712  0.415786556  0.437009864  0.461951799
## svi         0.3468674177  0.330415198  0.309283880  0.281608260  0.245177911
## gleason      0.1436779613  0.146778188  0.150949425  0.156678907  0.164800413
##
## (Intercept)      .              .              .
## lweight      -0.027603986 -0.064680201 -0.113137304
```

```

## age          0.017480107  0.018643148  0.020098181
## lbph         -0.040241264 -0.047425776 -0.056962692
## lcp          0.347616547  0.377657417  0.418431830
## pgg45        -0.005050263 -0.006739814 -0.009116838
## lpsa         0.491849702  0.528596455  0.575318051
## svi          0.196427346  0.129711598  0.035342349
## gleason      0.176722769  0.194999807  0.224585243

## functions to compute testing error with glmnet
error <- function(dat, glmnet, lam, form, loss=L2_loss) {
  x_inp <- model.matrix(form, data=dat)
  y_out <- dat$lcavol
  y_hat <- predict(glmnet, newx=x_inp, s=lam) ## see predict.elnet
  mean(loss(y_out, y_hat))
}

## training error at lambda=0.01
error(prostate_train, glmnet, lam=0.01, form=form)

## [1] 0.4385064

## testing error at lambda=0.01
error(prostate_test, glmnet, lam=0.01, form=form)

## [1] 0.5047688

## training error at lambda=0.1
error(prostate_train, glmnet, lam=0.1, form=form)

## [1] 0.4486907

## testing error at lambda=0.1
error(prostate_test, glmnet, lam=0.1, form=form)

## [1] 0.4914336

## training error at lambda=0.2
error(prostate_train, glmnet, lam=0.2, form=form)

## [1] 0.4653812

## testing error at lambda=0.2
error(prostate_test, glmnet, lam=0.2, form=form)

## [1] 0.4945718

## training error at lambda=0.3
error(prostate_train, glmnet, lam=0.3, form=form)

## [1] 0.481922

```

```
## testing error at lambda=0.3
error(prostate_test, glmnet, lam=0.3, form=form)
```

```
## [1] 0.5015496
```

```
## training error at lambda=0.4
error(prostate_train, glmnet, lam=0.4, form=form)
```

```
## [1] 0.4973256
```

```
## testing error at lambda=0.4
error(prostate_test, glmnet, lam=0.4, form=form)
```

```
## [1] 0.5088406
```

```
## training error at lambda=0.5
error(prostate_train, glmnet, lam=0.5, form=form)
```

```
## [1] 0.5115869
```

```
## testing error at lambda=0.5
error(prostate_test, glmnet, lam=0.5, form=form)
```

```
## [1] 0.5156281
```

```
## training error at lambda=0.6
error(prostate_train, glmnet, lam=0.6, form=form)
```

```
## [1] 0.524921
```

```
## testing error at lambda=0.6
error(prostate_test, glmnet, lam=0.6, form=form)
```

```
## [1] 0.5218096
```

```
## training error at lambda=0.7
error(prostate_train, glmnet, lam=0.7, form=form)
```

```
## [1] 0.5374869
```


```
## testing error at lambda=0.7
error(prostate_test, glmnet, lam=0.7, form=form)
```

```
## [1] 0.5274336
```

```
##Question 7:
```

```
err_train_1 <- sapply(glmnet$lambda, function(lam)
  error(prostate_train, glmnet, lam, form))
err_test_1 <- sapply(glmnet$lambda, function(lam)
  error(prostate_test, glmnet, lam, form))

## plot test/train error
plot(x=range(glmnet$lambda),
     y=range(c(err_train_1, err_test_1)),
     xlim=rev(range(glmnet$lambda)),
     type='n',
     xlab=expression(lambda),
     ylab='train/test error')
points(glmnet$lambda, err_train_1, pch=19, type='b', col='darkblue')
points(glmnet$lambda, err_test_1, pch=19, type='b', col='darkred')
legend('topright', c('train', 'test'), lty=1, pch=19,
      col=c('darkblue', 'darkred'), bty='n')
```



```
colnames(glmnet$beta) <- paste('lam =', glmnet$lambda)
print(glmnet$beta %>% as.matrix)
```

```
##          lam = 0.8    lam = 0.75    lam = 0.7    lam = 0.65    lam = 0.6
## (Intercept) 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
## lweight    0.126233723 0.123269298 0.119863278 0.1158716819 0.1111931576
## age        0.011482623 0.011670731 0.011877123 0.0121010799 0.0123455384
## lbph       -0.007449919 -0.008367575 -0.009363254 -0.0104653587 -0.0116905450
## lcp         0.207036526 0.211463094 0.216338206 0.2216006088 0.2273130852
## pgg45       0.001473672 0.001329880 0.001167084 0.0009827841 0.0007732439
## lpsa        0.290917892 0.298034308 0.305629735 0.3138070249 0.3226494596
## svi         0.394617510 0.393975652 0.392543932 0.3903246504 0.3871479640
## gleason     0.132832538 0.133493765 0.134175184 0.1349324324 0.1357928799
##          lam = 0.55    lam = 0.5    lam = 0.45    lam = 0.4
```

```
## (Intercept) 0.0000000000 0.0000000000 0.000000e+00 0.0000000000
## lweight 0.1057367423 0.0992845548 9.167569e-02 0.0826683974
## age 0.0126162948 0.0129137196 1.324527e-02 0.0136179442
## lbph -0.0130587545 -0.0145982826 -1.633988e-02 -0.0183250234
## lcp 0.2336122365 0.2404842462 2.481056e-01 0.2566365841
## pgg45 0.0005320808 0.0002563424 -6.281817e-05 -0.0004352185
## lpsa 0.3322362737 0.3427336970 3.542772e-01 0.3670625821
## svi 0.3827345083 0.3769402232 3.693573e-01 0.3595330231
## gleason 0.1368039677 0.1380061651 1.394762e-01 0.1413173664
## lam = 0.35 lam = 0.3 lam = 0.25 lam = 0.2 lam = 0.15
## (Intercept) 0.0000000000 0.0000000000 0.0000000000 0.0000000000 0.0000000000
## lweight 0.0719539790 0.059127616 0.043652593 0.024777969 0.001504802
## age 0.0140406458 0.014526957 0.015088490 0.015748487 0.016532948
## lbph -0.0206075810 -0.023258103 -0.026377963 -0.030098852 -0.034621150
## lcp 0.2662889676 0.277447149 0.290342311 0.305728439 0.324372008
## pgg45 -0.0008738898 -0.001398912 -0.002031353 -0.002810371 -0.003788173
## lpsa 0.3813402190 0.397429712 0.415786556 0.437009864 0.461951799
## svi 0.3468674177 0.330415198 0.309283880 0.281608260 0.245177911
## gleason 0.1436779613 0.146778188 0.150949425 0.156678907 0.164800413
## lam = 0.1 lam = 0.05 lam = 0
## (Intercept) 0.0000000000 0.0000000000 0.0000000000
## lweight -0.027603986 -0.064680201 -0.113137304
## age 0.017480107 0.018643148 0.020098181
## lbph -0.040241264 -0.047425776 -0.056962692
## lcp 0.347616547 0.377657417 0.418431830
## pgg45 -0.005050263 -0.006739814 -0.009116838
## lpsa 0.491849702 0.528596455 0.575318051
## svi 0.196427346 0.129711598 0.035342349
## gleason 0.176722769 0.194999807 0.224585243
```

##Question 8:

```
plot(x=range(glmnet$lambda),
     y=range(as.matrix(glmnet$beta)),
     type='n',
     xlab=expression(lambda),
     ylab='Coefficients')
for(i in 1:nrow(glmnet$beta)) {
  points(x=glmnet$lambda, y=glmnet$beta[i,], pch=19, col='#00000055')
  lines(x=glmnet$lambda, y=glmnet$beta[i,], col='#00000055')
}
text(x=0, y=glmnet$beta[,ncol(glmnet$beta)],
     labels=rownames(glmnet$beta),
     xpd=NA, pos=4, srt=45)
abline(h=0, lty=3, lwd=2)
```